

Team Members: Michas Szacillo (mas744), Alexandra Ward (amw349), Brandon Sweezer(bas329), and Abiy Tibebu (agt33)

Project 3 Proposal

Baseline System:

For the baseline, we decided to use a VERY simple Random Guess model. We decided to start off with a basic approach and randomly return a noun phrase answer for each question per paragraph. This will help us get an intuition of the QA problem and gain insight into the better methods we should implement on our final model which we describe below.

The noun phrases were parsed by using the TextBlob library, which looks at each given paragraph and returns a list of noun phrase tokens. We then randomly generate a number based on the length of the list, and return this as the answer. Obviously this results in a terrible F1 score because of the random nature and the small noun phrases that might not entirely answer the question. However, it serves as a basis for our research and development.

Baseline Results:

Exact Match	F1-Score
0.3595080416272469	0.542103817244337

Proposal:

To begin, we plan on trying some more simple models such as the Window Sliding model used in the critique article. We also plan on playing around with tagging systems to make picking out the answers in the paragraph easier (for example POS tags, or IOG tags). As we continue to improve our model, we will look into possible linguistic NP filters as well, since we know that most of the answers to the questions are phrases and not just words.

Similarly, we plan on classifying question types and semantic class types so that we can narrow down the possible answers even further to a corresponding return type (e.g. Person, Place, etc). We would do this by searching questions for identifying words such as “who” and “where” which would indicate the question is looking for a person or a place, respectively.

This answer classification could be improved from hard-coded rules by using a machine learning approach. We could train a neural network to identify the most likely NER tag which a question is looking for using a training corpus found online, or worst case, hand generated by us. Using this trained neural network, we could give it our questions and then (hopefully) get

the correct NER tag to look for in the sentence. Done correctly, this would vastly improve the answer parsing process, for we would have an accurate and reliable way to determine what sort of thing we are looking for in the paragraph. Once we know this, it should be simple to identify which of the possible answers in the paragraph is the right one using the following method.

After we have found an answer, we plan to evaluate the validity of that answer using a method we will call Statement Rearrangement Perplexity (SRP). SRP will rearrange the phrasing of the question and the previously identified answer possibilities so that they form a statement. Then it will evaluate the perplexity of each of these three sentences, using N-gram (we will not use anything higher than trigrams as the answers will not be very long) models generated on the paragraph and question. The least perplex answer will be selected as our final answer. We believe this will improve our accuracy, for the correct answer will have extremely low perplexity, as the rearranged statement will have almost identical phrasing to that of the answer in the paragraph.

Group Contribution:

Brandon: Wrote writeup, came up with some ideas for part 2, setup git repo and slack channel.

Alexandra: Contributed to writeup and ideas for proposal.

Michas: Contributed to the writeup, and wrote the random generation baseline.

Abiy: Added later to group, contributed to proposal ideas. Will contribute more in future.